

# Part III

# Synchronization

## Monitors

*That's been one of my mantras - focus and simplicity.  
Simple can be harder than complex:  
You have to work hard to get your thinking clean to make it simple.  
But it's worth it in the end because once you get there, you can move mountains.*


# ***What Is a Monitor? - Basics***

- Monitor is a highly structured programming language construct. It consists of
  - ❖ **private** variables and **private** procedures that can only be used within a monitor.
  - ❖ **constructors** that initialize the monitor.
  - ❖ A number of (public) **monitor procedures** that are available to the users.
- Note that monitors **have no public** data.
- A monitor is a mini-OS with monitor procedures as system calls.

# **Monitor: Mutual Exclusion 1/2**

- **No more than one process** can be executing *in* a monitor. Thus, **mutual exclusion** is automatically guaranteed in a monitor.
- When a process calls a monitor procedure and enters the monitor successfully, it is the **only** process **executing** in the monitor.
- When a process calls a monitor procedure and the monitor has a process executing, the caller is blocked **outside of the monitor**.

# Monitor: Mutual Exclusion 2/2



- If there is a process executing in a monitor, any process that calls a monitor procedure is blocked **outside** of the monitor.
- When the monitor has no executing process, one process will be let in.

# Monitor: Syntax

```
monitor Monitor-Name
{
    local variable declarations;

    Procedure1 (...)
    { // statements };
    Procedure2 (...)
    { // statements };
    // other procedures
    {
        // initialization
    }
}
```

- All variables are **private**.  
**Why? Exercise!**
- **Monitor procedures are public**; however, some procedures may be private so that they can only be used within a monitor.
- **Initialization procedures** (i.e., **constructors**) execute only once when the monitor is created.

# Monitor: A Very Simple Example

```
monitor IncDec
```

```
{
```

```
    int count;
```

```
    void Increase(void)
    { count++; }
```

```
    void Decrease(void)
    { count--; }
```

```
    int GetData(void)
    { return count; }
```

```
    { count = 0; }
```

```
}
```

```
process Increment
```

```
while (1) {
```

```
    // do something
```

```
    IncDec.Increase();
```

```
    cout <<
```

```
        IncDec.GetData();
```

```
    // do something
```

```
}
```

Is the printed value the one just updated?

initialization

# ***Condition Variables***

- Mutual exclusion is an easy task with monitors.
- While a process is executing **in** a monitor, it may have to wait until an event occurs.
- Each programmer-defined event is conceptually represented by a ***condition variable***.
- A condition variable, or a condition, has a private waiting list, and two public methods: `signal` and `wait`.
- Note that a condition variable **has no value** and **cannot be modified**.

# **Condition wait**

- Let `cv` be a condition variable. The use of methods `signal` and `wait` on `cv` are `cv.signal()` and `cv.wait()`.
- Condition wait and condition signal can only be used ***in a monitor***.
- A process that executes a condition wait **blocks immediately** and is put into the waiting list of that condition variable. **The monitor becomes “empty” (i.e., no executing process inside).**
- This means that this process is waiting for the indicated event to occur.



# **Condition** *signal*

- Condition `signal` is used to indicate an event has occurred.
- If there are processes waiting on the signaled condition variable, **one of them** will be released.
- If there is **no waiting process** waiting on the signaled condition variable, **this signal is lost as if it never happens**.
- **Consider the released process (from the signaled condition) and the process that signals. There are **two** processes executing in the monitor, and mutual exclusion is violated! Something has to be done to fix this problem.**

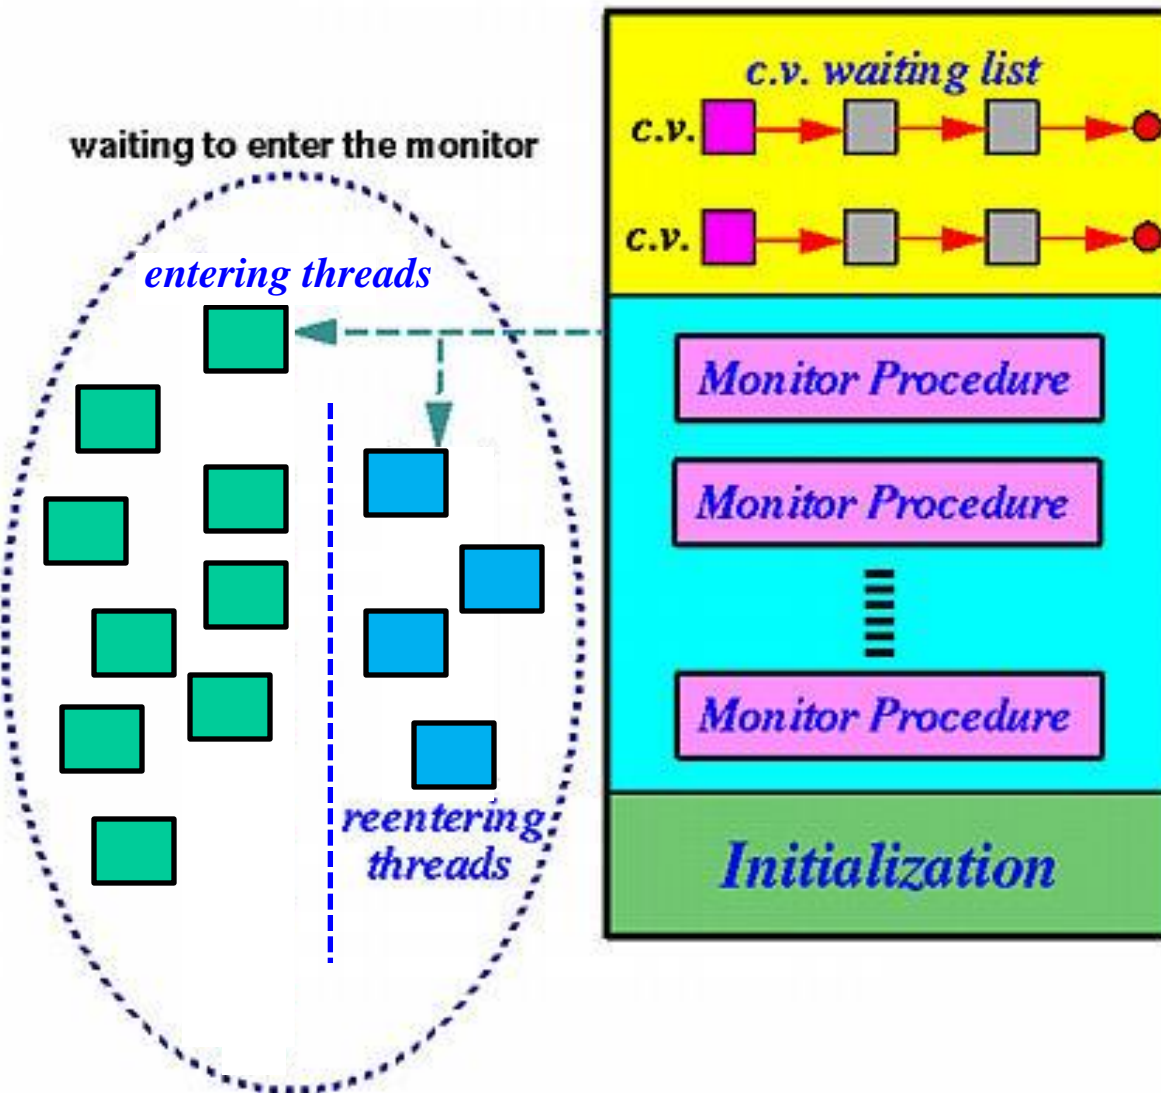
# ***Two Types of Monitors***

- After a signal, the released process and the signaling process may be executing in the monitor.
- There are **two** approaches to address this issue:
  - ❖ **Hoare Type** (proposed by C. A. R. Hoare):  
The released process takes the monitor and the signaling process **waits somewhere**.
  - ❖ **Mesa Type** (proposed by Lampson and Redell): The released process **waits somewhere** and the signaling process continues to use the monitor. This is also used in Java.

# ***What Do You Mean by “Waiting Somewhere”?***

- The signaling process (Hoare type) or the released process (Mesa type) must **wait somewhere**.
- You could consider there is a **waiting bench** for these processes to wait.
- Hence, each process that involves in a monitor call may be in one of the four states:
  - ❖ **Active**: The running one.
  - ❖ **Entering**: Those blocked by the monitor.
  - ❖ **Waiting**: Those waiting on a condition variable.
  - ❖ **Inactive**: Those waiting on the waiting bench.

# Monitor with Condition Variables



- Processes blocked due to signal/wait are in the **re-entry** list (*i.e.*, waiting bench).
- When the monitor is free, a process is released from either **entry** or **re-entry**.

# What Is the Major Difference?

Condition    UntilHappen;

// Hoare Type

```
if (!event)
    UntilHappen.wait();
```

// Mesa Type


```
while (!event)
    UntilHappen.wait();
```

***Unless stated otherwise,  
we only use the Hoare  
type monitors in this  
course.***

With **Hoare** type, once a signal arrives, the signaler yields the monitor to the released process and the condition is not changed. Thus, an **if** is sufficient.

With **Mesa** type, the released process may be waiting for a while before it runs. During this period, other processes may be in the monitor and change the condition. It is better to check the condition again with a **while**!

# Monitor: Producer/Consumer



```
monitor ProdCons
{
    int count, in, out;
    int Buf[SIZE];
    condition
        UntilFull,
        UntilEmpty;

    procedure PUT(int);
    procedure GET(int *);
    { count = 0}
}
```

# Monitor: *PUT()* and *GET()*

```
void PUT(int X)
{
    if (count == SIZE)
        UntilEmpty.wait();
    Buf[in] = X;
    in = (in+1)%SIZE;
    count++;
    if (count == 1)
        UntilFull.signal();
}
```

```
void GET(int *X)
{
    if (count == 0)
        UntilFull.wait();
    *X = Buf[out];
    out=(out+1)%SIZE;
    count--;
    if (count == SIZE-1)
        UntilEmpty.signal();
}
```

# Run This Solution with Mesa?

**Buffer Size = 2**

$P_1$	$P_2$	$P_3$	$P_4$	$C_1$	Count
					0
Add 1 item					1
	Add 1 item				2
		Wait UntilEmpty			2
				Take 1 item	1
				Signal UntilEmpty	1
			Add 1 item		2
By the time P3 finally runs, there is no space!					2
		No space!			2

if it is a Hoare monitor,  
 $P_3$  runs immediately

Under Mesa,  $C_1$  continues.

Upon exit, the monitor becomes empty  
and selects  $P_4$  to enter



# ***Dining Philosophers: Again!***

- In addition to **thinking** and **eating**, a philosopher has one more state, **hungry**, in which he is trying to get chopsticks.
- We use an array `state[]` to keep track the state of a philosopher. Thus, philosopher `i` can eat (*i.e.*, `state[i] = EATING`) only if his neighbors are not eating (*i.e.*, `state[(i+4)%5]` and `state[(i+1)%5]` are not `EATING`).

# ***Monitor Definition***

```
monitor philosopher
{
    enum { THINKING, HUNGRY,
           EATING } state[5];
    condition self[5];
    private: CanEat(int);

    procedure GET(int);
    procedure PUT(int);

    { for (i=0; i<5; i++)
        state[i] = THINKING;
    }
}
```

# ***The CanEat() Procedure***

```
void CanEat(int k)
{
    if ( (state[(k+4)%5] != EATING) &&
         (state[k] == HUNGRY) &&
         (state[(k+1)%5] != EATING) ) {
        state[k] = EATING;
        self[k].signal();
    }
}
```

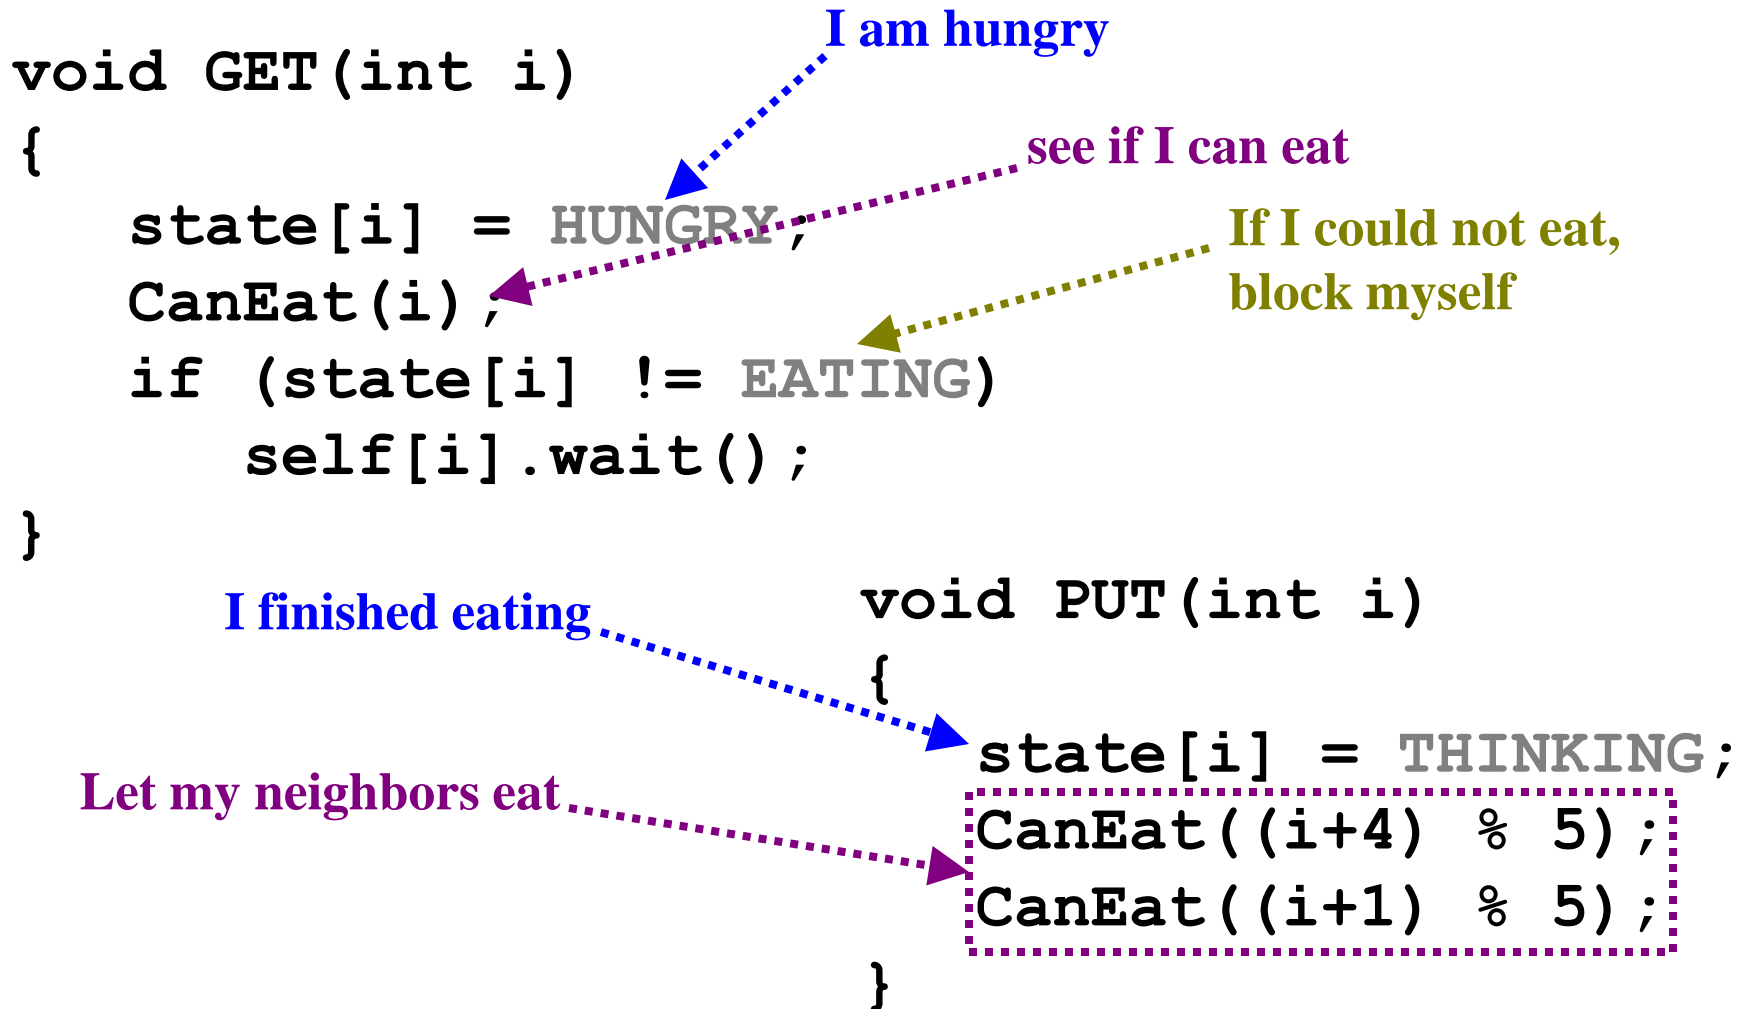
the left and right neighbors of philosopher  $k$  are not eating

philosopher  $k$  is hungry

The diagram illustrates the logic of the CanEat() procedure. It shows the code with annotations and visual cues: blue dashed boxes and arrows group the conditions for neighbors not eating and k being hungry, while an orange dashed box and arrow highlight the condition that k is hungry.

- If the left and right neighbors of philosopher  $k$  are **not eating** and philosopher  $k$  is **hungry**, then philosopher  $k$  can **eat**. Thus, release him!

# **The *GET()* and *PUT()* Procedures**



**Which type of monitor am I using?**

# How about Deadlock?

```
void CanEat(int k)
{
    if ((state[(k+4)%5] != EATING) &&
        (state[k] == HUNGRY) &&
        (state[(k+1)%5] != EATING)) {
        state[k] = EATING;
        self[k].signal();
    }
}
```

- This solution does not have deadlock, because
  1. The only place where eating permission is granted is in procedure `CanEat()`, and
  2. Philosopher *k* can eat only if he could get both chopsticks (i.e., no hold and wait and no circular waiting).

# How about Bounded Waiting?

```
void CanEat(int k)
{
    if ((state[(k+4)%5] != EATING) &&
        (state[k] == HUNGRY) &&
        (state[(k+1)%5] != EATING)) {
        state[k] = EATING;
        self[k].signal();
    }
}
```

- **Question:** The Progress condition is met and could be proved easily. How about the **Bounded Waiting** condition? More precisely, is it possible that some philosophers can continue the process of thinking and eating and block some others indefinitely? **Exercise.**

# ***The Reader-Writer Problem: Again!***

- Let us add one minor modification to the original reader-writer problem to make it a bit more realistic:
  - If a writer is waiting, the new readers should wait their turn, even though it is safe to proceed if there are readers reading.

# ***Monitor Definition***

```
monitor reader-writer
{
    int reading = 0;    // reading readers
    int writing = 0;     // writing writers
    int writers = 0;    // waiting writers

    condition  Take_Turn;

    procedure read_REQUEST(void) ;
    procedure read_RELEASE(void) ;
    procedure write_REQUEST(void) ;
    procedure write_RELEASE(void) ;
}
```



# ***Readers and Writers***

## **Reader**

```
while (1)
{
    // do something
    read_REQUEST();
    // reading
    read_RELEASE();
    // do something
}
```

## **Writer**

```
while (1)
{
    // do something
    write_REQUEST();
    // writing
    write_RELEASE();
    // do something
}
```

# ***Monitor Code for Readers***

```
void read_REQUEST()
{
    if (writers > 0)        // if there are writers waiting,
        Take_Turn.Wait();  //      wait until released
    if (writing > 0)        // if there is a writer writing,
        Take_Turn.Wait();  //      wait, of course
    reading++;              // because no writer is writing,
                            //      a reader can read
}

void read_RELEASE()
{
    reading--;              // a reader has done reading
    Take_Turn.Signal();    // let one reader/writer to go
}
```

# Monitor Code for Writers

```
void write_REQUEST()
```

this means reading or writing being non-zero

```
{  
    writers++;           // one more write request  
    if (reading || writing) // if there is readers reading  
        Take_Turn.Wait(); // or w writer writing, wait  
    writing++;           // if no readers reading and no  
}                       // writer writing, then go!
```

```
void write_RELEASE()
```

```
{  
    writing--;           // reduce writing count  
    writers--;         // reduce writer count  
    Take_Turn.Signal(); // let some one to proceed  
}
```

# ***Hoare Type vs. Mesa Type***

- When a signal occurs, **Hoare** type monitor uses **two** context switches, one switching the signaling process out and the other switching the released in. However, **Mesa** type monitor uses **one**.
- Process scheduling must be very **reliable** with **Hoare** type monitors to ensure once the signaling process is switched out the next one must be the released process. **Why?**
- With **Mesa** type monitors, a condition may be evaluated multiple times. However, **incorrect signals** will do less harm because every process checks its own condition.

# ***Semaphores vs. Monitors***

<b>Semaphores</b>	<b>Monitors</b>
<b>Can be used anywhere, but should not be in a monitor</b>	<b>Can only be accessed with monitor procedure calls</b>
<b>No connection between the semaphore and the data this semaphore protects</b>	<b>Data and access procedures are in the same place (i.e., a monitor)</b>
<b>Semaphores are low level assembly language-like instructions</b>	<b>Monitors are well-structured higher-level construct</b>
<b>Not easy to use and prone to bugs</b>	<b>Easy of use and good protection of vital data</b>

# Semaphores vs. Conditions

Semaphores	Condition Variables
Can be used anywhere, but not in a monitor	Can only be used in monitors
<code>wait()</code> does not always block its caller	<code>wait()</code> <b>always</b> blocks its caller
<code>signal()</code> either releases a process, or increases the semaphore counter	<code>signal()</code> either releases a process, or the signal is <b>lost</b> as if it never occurs
If <code>signal()</code> releases a process, the caller and the released <b><i>both continue</i></b>	If <code>signal()</code> releases a process, either the caller or the released continues, but <b><i>not both</i></b>

# ***Semaphore and Monitor Equivalence***

- In terms of expressive power, semaphores and monitors are equivalent.
- A semaphore can be implemented with a monitor. This is easy and is your homework.
- Conversely, a monitor and its condition variables may also be simulated with multiple semaphore, although this is tedious. See weekly reading list.
- Therefore, semaphores and monitors are equivalent because one may be implemented by the other.

# Monitors with ***ThreadMentor***



# ***Monitor: Definition***

```
class MyMon::public Monitor
{
    public:
        MyMon(); // constructor
        MonitorProcedure-1();
        MonitorProcedure-2();
        // other procedures
    private:
        // variables used in
        // this monitor
};
```

- A monitor must be a derived class of class `Monitor`.
- The initialization part should be in constructors.
- Make monitor procedures `public`.
- Local variables should be `private/protected`.

# Monitor: Monitor Procedures

```
int MyMon::MonProc (...)  
{  
    MonitorBegin();  
    // other statements  
    // of this procedure  
    MonitorEnd();  
}
```

*MonitorBegin() locks the monitor and MonitorEnd() unlocks it. Thus, mutual exclusion is guaranteed.*

- Monitor procedures are C/C++ functions.
- Before you do anything, call `MonitorBegin()`.
- Before exit, call `MonitorEnd()`.
- The following is **wrong**:

```
int MyMon::MonProc ()  
{  
    MonitorBegin();  
    // other stuffs  
    return 0;  
    MonitorEnd();  
}
```

# Monitor: A Simple Example

```
Class Count
    ::public Monitor
{
    public:
        int  Inc() ;
        int  Dec() ;
        void Count() ;
    private:
        int  Counter;
}

Count::Count(void)
{ Counter = 0; }

int Count::Inc()
{
    MonitorBegin() ;
    Counter++;
    MonitorEnd() ;
    return Count;
}

int Count::Dec()
{
    MonitorBegin() ;
    Count--;
    MonitorEnd() ;
    return Count;
}
```

# ***Monitor: Condition Variables***

**Condition**    **Event;**

**Event.Wait();**

**Event.Signal();**

- **Condition** is a class and has two methods, **Wait()** and **Signal()**.
- **Waiting on a condition variable means waiting for that event to occur.**
- **Signaling a condition variable means that the event has occurred.**

# Philosopher Monitor Definition

condition variable pointers, one for each philosopher

```
class Mon::public Monitor
{
    public:
        Mon();
        GET(int); PUT(int);
    private:
        Condition *Self[5];
        int State[5];
        int CanEat(int);
};
```

get and put  
chopsticks

GET(int); PUT(int);

Condition \*Self[5];

int State[5];

int CanEat(int);

```
Mon::Mon()
```

```
{ int i;
```

```
for (i=0; i < 5; i++){
```

```
    State[i] = THINKING;
```

```
}
```

are both chopsticks available?

state of each philosopher

# Philosopher Monitor Implementation

```
int Mon::CanEat(int k)
{
    if ((state[(k+4)%5] !=
        EATING)
        &&(state[k] ==
        HUNGRY)
        && (state[(k+1)%5] !=
        EATING)) {
        state[k] = EATING;
        self[k].signal();
    }
}
```

check to see if I can eat

```
void Mon::GET(int k)
{
    MonitorBegin();
    state[k] = HUNGRY;
    CanEat(k);
    if (state[k] != EATING)
        self[k].wait();
    MonitorEnd();
}
```

if I cannot eat, wait

# ***Specifying a Monitor Type***

```
MyMonitor::MyMonitor(char *Name)
    : Monitor(Name, HOARE)
{
    // initialization here
}
```

Replace HOARE with MESA  
if you wish to use a Mesa  
type monitor.

- A monitor type must be specified in your monitor constructor.
- Use HOARE or MESA for Hoare type and Mesa type monitors.

**The End**