

# Part III

# Synchronization

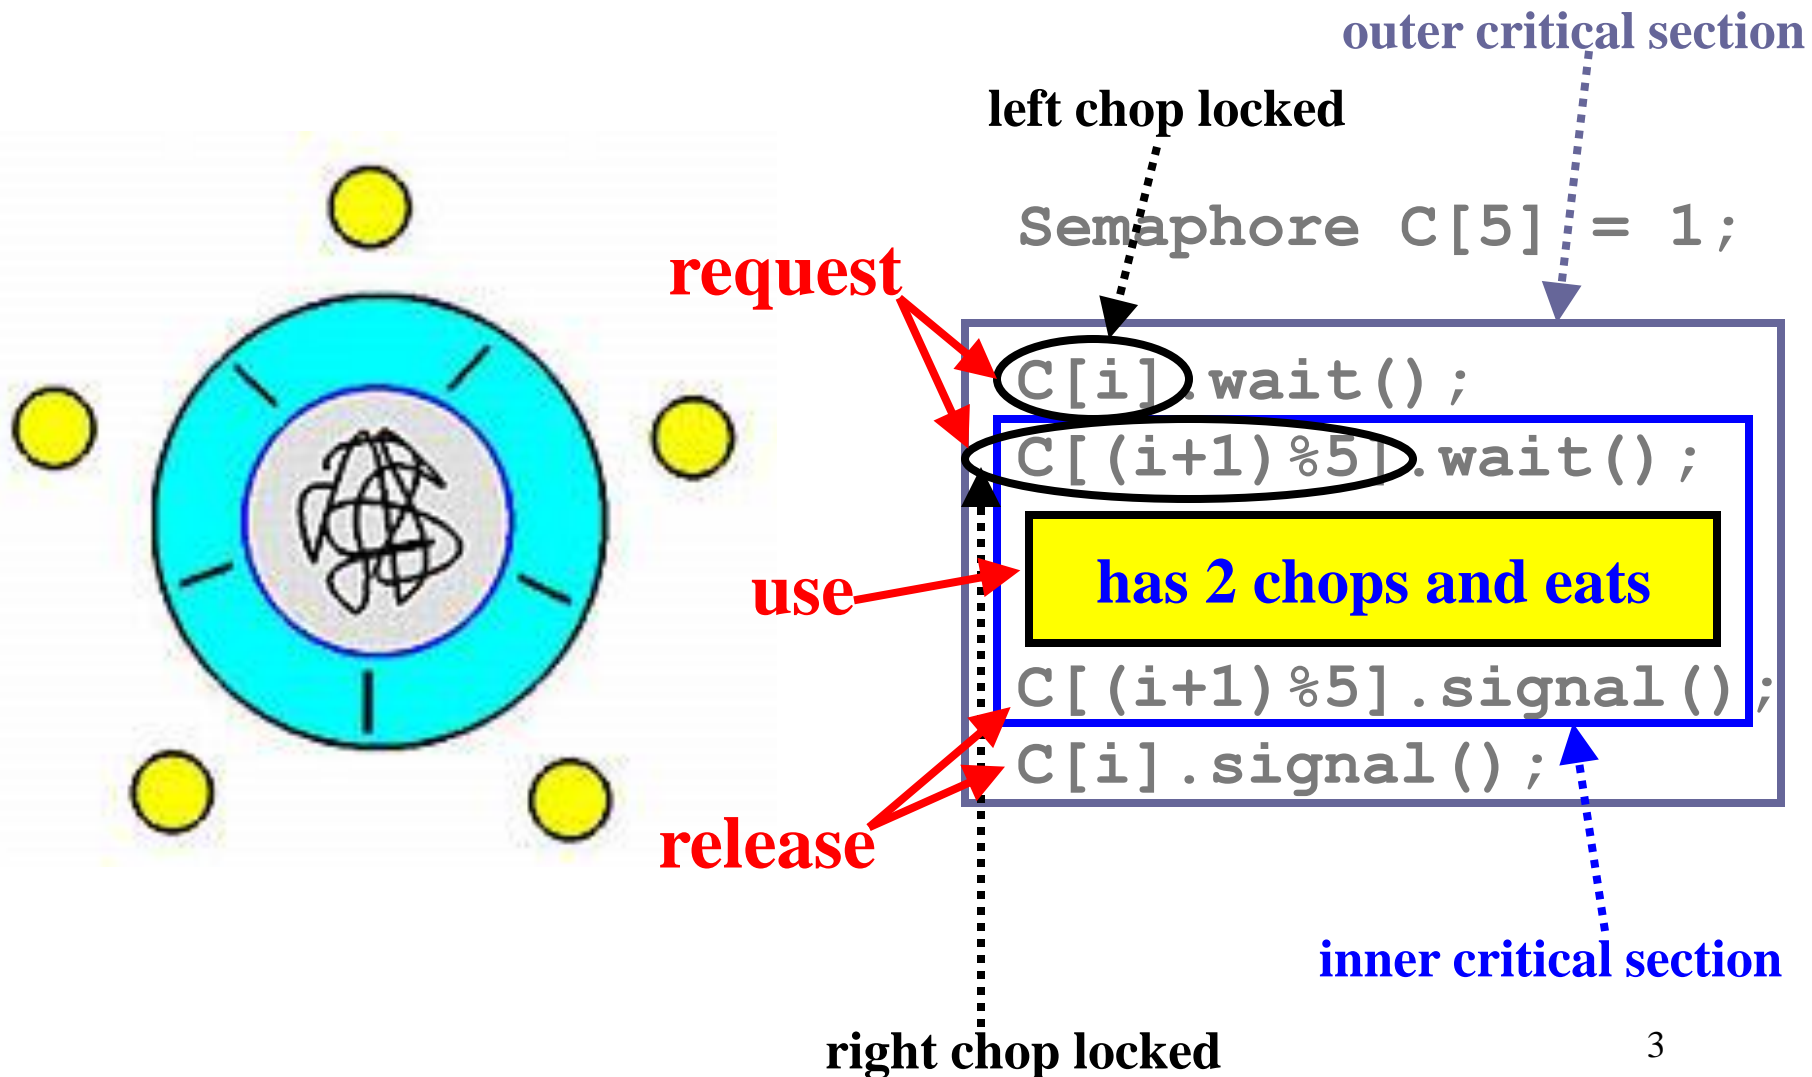
## Deadlocks and Livelocks

*You think you know when you learn,  
are more sure when you can write,  
even more when you can teach,  
but certain when you can program.*

# ***System Model: 1/2***

- System resources are used in the following way:
  - ❖ **Request**: If a process makes a request (i.e., semaphore wait or monitor acquire) to use a system resource which cannot be granted immediately, then the requesting process blocks until it can acquire the resource successfully.
  - ❖ **Use**: The process operates on the resource (i.e., in critical section).
  - ❖ **Release**: The process releases the resource (i.e., semaphore signal or monitor release).

# System Model: 2/2



# ***Deadlock: Definition***

- A set of processes is in a ***deadlock*** state when every process in the set is waiting for an event that can only be caused by another process in the same set.
- The key here is that processes are all in the ***waiting state***.

# ***Deadlock Necessary Conditions***

- ***If a deadlock occurs, then each of the following four conditions must hold.***
  - ❖ ***Mutual Exclusion***: At least one resource must be held in a non-sharable way.
  - ❖ ***Hold and Wait***: A process must be holding a resource and waiting for another.
  - ❖ ***No Preemption***: Resource cannot be preempted.
  - ❖ ***Circular Waiting***:  $P_1$  waits for  $P_2$ ,  $P_2$  waits for  $P_3$ , ...,  $P_{n-1}$  waits for  $P_n$ , and  $P_n$  waits for  $P_1$ .

# **Deadlock Necessary Conditions**

- **Note that the conditions are necessary.**
- **This means if a deadlock occurs **ALL** conditions are met.**
- **Because  $p \Rightarrow q$  is equivalent to  $\neg q \Rightarrow \neg p$ , where  $\neg q$  means not all conditions are met and  $\neg p$  means no deadlock, as long as one of the four conditions fails there will be no deadlock.**


# ***Deadlock Prevention: 1/7***

- ***Deadlock Prevention*** means making sure deadlocks never occur.
- To this end, if we are able to make sure at least one of the four conditions fails, there will be no deadlock.

# **Deadlock Prevention: 2/7**

## **Mutual Exclusion**

- **Mutual Exclusion:** Some sharable resources must be accessed exclusively, which means we cannot deny the mutual exclusion




The use of these five chopsticks must be mutually exclusive



# ***Deadlock Prevention: 3/7***

## ***Hold and Wait***

- ***Hold and Wait***: A process holds some resources and requests for other resources.




# ***Deadlock Prevention: 4/7***

## ***Hold and Wait***


- ***Solution***: Make sure no process can hold some resources and then request for other resources.
- Two strategies are possible (***the monitor solution to the philosophers problem***):
  - ❖ A process must acquire ***all*** resources before it runs.
  - ❖ When a process requests for resources, it must hold none (i.e., returning resources before requesting for more).
- ***Resource utilization*** may be low, since many resources will be held and unused for a long time.
- ***Starvation*** is possible. A process that needs some popular resources may have to wait indefinitely.

# ***Deadlock Prevention: 5/7***

## ***Hold and Wait***



If weirdo is faster than #1, #1 cannot eat  
and the weirdo or #4 can eat but not both.  
If weirdo is slower than #1, #4 can eat  
Since there is no hold and wait,  
there is no deadlock.



In this case, #4 has no right neighbor  
and can take his right chop.  
Since there is no hold and wait,  
there is no deadlock.

The monitor solution with THINKING-HUNGRY-EATING states forces a philosopher to have both chops before eating. Hence, no hold-and-wait.

# ***Deadlock Prevention: 6/7***


## ***No Preemption***

- This means resources being held by a process cannot be taken away (i.e., no preemption).
- To negate this no preemption condition, a process may deallocate all resources it holds so that the other processes can use.
- This is sometimes not doable. For example, while philosopher *i* is eating, his neighbors cannot take *i*'s chops away forcing *i* to stop eating.
- Moreover, some resources cannot be reproduced cheaply (e.g., printer).

# **Deadlock Prevention: 7/7**

## **Circular Waiting**

- **Circular Waiting:**  $P_1$  waits for  $P_2$ ,  $P_2$  waits for  $P_3$ , ...,  $P_{n-1}$  waits for  $P_n$ , and  $P_n$  waits for  $P_1$ .



The weirdo, 4-chair, and monitor solutions all avoid circular waiting and there is no deadlock.

Resources can be ordered in a hierarchical way.

A process can only acquire resources higher than those it has.

To acquire lower order resources a process must release all resources higher than or equal to that of the acquiring one.

As a result, no deadlock can happen.

**Prove this yourself.**

# ***Livelock: 1/3***

- ***Livelock***: If two or more processes continually repeat the same interaction in response to changes in the other processes without doing any useful work.
- These processes are ***not*** in the waiting state, and they are running concurrently.
- This is different from a deadlock because in a deadlock all processes are in the waiting state.

# ***Livelock: 2/3***

```
MutexLock Mutex1, Mutex2;

Mutex1.Lock();           // lock Mutex1
while (Mutex2.isLocked()) { // loop until Mutex2 is open
    Mutex1.Unlock();      // release Mutex1 (yield)
    // wait for a while   // wait for a while
    Mutex1.Lock();        // reacquire Mutex1
}                          // OK, Mutex2 is open
Mutex2.Lock();           // lock Mutex2. have both

Mutex2.Lock();
while (Mutex1.isLocked()) {
    Mutex2.Unlock();
    // wait for a while
    Mutex2.Lock();
}
Mutex1.Lock();
```

**Both processes try to acquire two locks and they yield to each other** 15

# ***Livelock: 3/3***

- Process 1 locks `Mutex1` first. If `Mutex2` is not locked, process 1 acquires it. Otherwise, process 1 yields `Mutex1`, waits for a while (for process 2 to take `Mutex1` and finish its task), reacquires `Mutex1`, and checks again `Mutex2` is open.
- Process 2 does this sequence the same way with the role of `Mutex1` and `Mutex2` switched.
- To avoid this type of livelock, ***order the locking sequence in a hierarchical way*** (i.e., both lock `Mutex1` ***first*** followed by `Mutex2`). Thus, only one process can lock both locks successfully.



# The End